

# **HEART DISEASE CLASSIFICATIONS USING MACHINE LEARNING**

**Project Report  
Submitted to**

**DEVI AHILYA VISHWAVIDHYALYA INDORE**



**In Partial Fulfillment of the Requirements  
for the Degree of**

**MASTER OF TECHNOLOGY**

in

**BIG DATA ANALYTIC**

by

**Manish Kumar Singh  
DS7B-2206**

**Under the Supervision of**

**Dr. Dinesh Bhati**  
Assistant Professor



**SCHOOL OF DATA SCIENCE AND  
FORECASTING**

**2022-23**

## **DECLARATION**

I hereby declare that the work presented in this project report entitled “HEART DISEASE CLASSIFICATIONS USING MACHINE LEARNING” was carried out by me. I have not submitted the matter embodied in this report for the award of any other degree or diploma of any other University or Institute. I have given due credit to the original authors/sources for all the words, ideas, diagrams, graphics, computer programs, experiments, results, that are not my original contribution. I have used quotation marks to identify verbatim sentences and given credit to the original authors/sources. I affirm that no portion of my work is plagiarized, and the experiments and results reported in the project report are not manipulated. In the event of a complaint of plagiarism and the manipulation of the experiments and results, I shall be fully responsible and answerable.

Manish Kumar Singh

DS7B-2206

(Candidate Signature)

## CERTIFICATE

Certified that Manish Kumar Singh (roll number DS7B-2206) has carried out the project work titled “Titanic decision tree algorithm” for the award of Master of Technology in Big Data Analytic DEVI AHILYA VISHWAVIDHYALYA INDORE under my supervision. The work embodies results of work, carried out by the student himself/herself and the contents of the work do not form the basis for the award of any other degree to the candidate or to anybody else from this or any other University/Institution.

**Name of Supervisor:** Dr. Dinesh Bhati

**Head of the Department:** Dr. V.B Gupta

**Designation:** Assistant Professor

**Signature**

**Signature**

**Date**

**Date**

## ABSTRACT

You have just been hired as a **Data Scientist** at a Hospital with an alarming number of patients coming in reporting various cardiac symptoms. A cardiologist measures vitals & hands you this data to **perform Data Analysis** and **predict** whether certain patients have Heart Disease. We would like to make a **Machine Learning algorithm** where we can train our AI to learn & improve from experience. Thus, we would want to **classify** patients as either positive or negative for Heart Disease.

<b>INDEX</b>
INTRODUCTION
Read Data
Data Exploration
Creating Dummy Variables
Creating Model for Logistic Regression
Normalize Data
Sigmoid Function
Forward and Backward Propagation
Cost Function
Gradient Descent
Sklearn Logistic Regression
K-Nearest Neighbour (KNN) Classification
Support Vector Machine (SVM) Algorithm
Naive Bayes Algorithm
Decision Tree Algorithm
Random Forest Classification
Comparing Models
Confusion Matrix

# INTRODUCTION

We have a data which classified if patients have heart disease or not according to features in it.  
We will try to use this data to create a model which tries predict if a patient has this disease or not.  
We will use logistic regression (classification) algorithm.

```
In [2]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load in

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split

# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list the files in the input directory

import os
print(os.listdir("../input"))

# Any results you write to the current directory are saved as output.
```

['All Users', 'Default', 'Default User', 'desktop.ini', 'PC', 'Public']

## Read Data

```
In [40]: # We are reading our data
df = pd.read_csv("heart.csv")
```

```
In [41]: # First 5 rows of our data
df.head()
```

Out[41]:

	age	sex	cp	trtbps	chol	fbs	restecg	thalach	exng	oldpeak	slope	caa	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

Data contains;

- age - age in years
- sex - (1 = male; 0 = female)


- cp - chest pain type
- trestbps - resting blood pressure (in mm Hg on admission to the hospital)
- chol - serum cholestoral in mg/dl
- fbs - (fasting blood sugar > 120 mg/dl) (1 = true; 0 = false)
- restecg - resting electrocardiographic results
- thalach - maximum heart rate achieved
- exang - exercise induced angina (1 = yes; 0 = no)
- oldpeak - ST depression induced by exercise relative to rest
- slope - the slope of the peak exercise ST segment
- ca - number of major vessels (0-3) colored by flourosopy
- thal - 3 = normal; 6 = fixed defect; 7 = reversable defect
- target - have disease or not (1=yes, 0=no)

## Data Exploration

```
In [19]: df.target.value_counts()
```

```
Out[19]: 1    165  
         0    138  
         Name: target, dtype: int64
```


```
In [20]: sns.countplot(x="target", data=df, palette="bwr")  
plt.show()
```



```
In [21]: countNoDisease = len(df[df.target == 0])
countHaveDisease = len(df[df.target == 1])
print("Percentage of Patients Haven't Heart Disease: {:.2f}%".format((countNoDisease / (len(df.target) * 100))))
print("Percentage of Patients Have Heart Disease: {:.2f}%".format((countHaveDisease / (len(df.target) * 100))))
```

Percentage of Patients Haven't Heart Disease: 45.54%  
Percentage of Patients Have Heart Disease: 54.46%

```
In [22]: sns.countplot(x='sex', data=df, palette="mako_r")
plt.xlabel("Sex (0 = female, 1= male)")
plt.show()
```



```
In [23]: countFemale = len(df[df.sex == 0])
countMale = len(df[df.sex == 1])
print("Percentage of Female Patients: {:.2f}%".format((countFemale / (len(df.sex) * 100))))
print("Percentage of Male Patients: {:.2f}%".format((countMale / (len(df.sex) * 100))))
```

Percentage of Female Patients: 31.68%  
Percentage of Male Patients: 68.32%




```
In [24]: df.groupby('target').mean()
```


```
Out[24]:
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	target
0	56.601449	0.826087	0.478261	134.398551	251.086957	0.159420	0.449275	139.101449	0.5
1	52.496970	0.563636	1.375758	129.303030	242.230303	0.139394	0.593939	158.466667	0.1


```
In [25]: pd.crosstab(df.age,df.target).plot(kind="bar",figsize=(20,6))
plt.title('Heart Disease Frequency for Ages')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.savefig('heartDiseaseAndAges.png')
plt.show()
```




```
In [26]: pd.crosstab(df.sex,df.target).plot(kind="bar",figsize=(15,6),color=['#1CA53B','#A52A2A'])
plt.title('Heart Disease Frequency for Sex')
plt.xlabel('Sex (0 = Female, 1 = Male)')
plt.xticks(rotation=0)
plt.legend(["Haven't Disease", "Have Disease"])
plt.ylabel('Frequency')
plt.show()
```




```
In [31]: plt.scatter(x=df.age[df.target==1], y=df.thalach[(df.target==1)], c="red")
plt.scatter(x=df.age[df.target==0], y=df.thalach[(df.target==0)])
plt.legend(["Disease", "Not Disease"])
plt.xlabel("Age")
plt.ylabel("Maximum Heart Rate")
plt.show()
```




```
In [36]: pd.crosstab(df.slope,df.target).plot(kind="bar",figsize=(15,6),color=['#DAF7A6', '#FF7F0E'],
plt.title('Heart Disease Frequency for Slope')
plt.xlabel('The Slope of The Peak Exercise ST Segment ')
plt.xticks(rotation = 0)
plt.ylabel('Frequency')
plt.show()
```



```
In [37]: pd.crosstab(df.fbs,df.target).plot(kind="bar",figsize=(15,6),color=['#FFC300', '#550154'],
plt.title('Heart Disease Frequency According To FBS')
plt.xlabel('FBS - (Fasting Blood Sugar > 120 mg/dl) (1 = true; 0 = false)')
plt.xticks(rotation = 0)
plt.legend(["Haven't Disease", "Have Disease"])
plt.ylabel('Frequency of Disease or Not')
plt.show()
```



```
In [38]: pd.crosstab(df.cp,df.target).plot(kind="bar",figsize=(15,6),color=['#11A5AA','#A511A5'],
plt.title('Heart Disease Frequency According To Chest Pain Type')
plt.xlabel('Chest Pain Type')
plt.xticks(rotation = 0)
plt.ylabel('Frequency of Disease or Not')
plt.show()
```



## Creating Dummy Variables

Since 'cp', 'thal' and 'slope' are categorical variables we'll turn them into dummy variables.

```
In [42]: a = pd.get_dummies(df['cp'], prefix = "cp")
b = pd.get_dummies(df['thal'], prefix = "thal")
c = pd.get_dummies(df['slope'], prefix = "slope")
```

```
In [43]: frames = [df, a, b, c]
df = pd.concat(frames, axis = 1)
df.head()
```

Out[43]:

	age	sex	cp	trtbps	chol	fbs	restecg	thalach	exng	oldpeak	...	cp_1	cp_2	cp_3	thal
0	63	1	3	145	233	1	0	150	0	2.3	...	0	0	1	
1	37	1	2	130	250	0	1	187	0	3.5	...	0	1	0	
2	41	0	1	130	204	0	0	172	0	1.4	...	1	0	0	
3	56	1	1	120	236	0	1	178	0	0.8	...	1	0	0	
4	57	0	0	120	354	0	1	163	1	0.6	...	0	0	0	

5 rows × 25 columns

```
In [44]: df = df.drop(columns = ['cp', 'thal', 'slope'])
df.head()
```

Out[44]:

	age	sex	trtbps	chol	fbs	restecg	thalach	exng	oldpeak	caa	...	cp_1	cp_2	cp_3	thal_
0	63	1	145	233	1	0	150	0	2.3	0	...	0	0	1	(
1	37	1	130	250	0	1	187	0	3.5	0	...	0	1	0	(
2	41	0	130	204	0	0	172	0	1.4	0	...	1	0	0	(
3	56	1	120	236	0	1	178	0	0.8	0	...	1	0	0	(
4	57	0	120	354	0	1	163	1	0.6	0	...	0	0	0	(

5 rows × 22 columns

## Creating Model for Logistic Regression

We can use sklearn library or we can write functions ourselves. Let's them both. Firstly we will write our functions after that we'll use sklearn library to calculate score.

```
In [45]: y = df.target.values
x_data = df.drop(['target'], axis = 1)
```

## Normalize Data

```
In [46]: # Normalize
x = (x_data - np.min(x_data)) / (np.max(x_data) - np.min(x_data)).values
```

C:\Users\PC\anaconda3\lib\site-packages\numpy\core\fromnumeric.py:84: FutureWarning: In a future version, DataFrame.min(axis=None) will return a scalar min over the entire DataFrame. To retain the old behavior, use 'frame.min(axis=0)' or just 'frame.min()'

```
    return reduction(axis=axis, out=out, **passkwargs)
```

C:\Users\PC\anaconda3\lib\site-packages\numpy\core\fromnumeric.py:84: FutureWarning: In a future version, DataFrame.max(axis=None) will return a scalar max over the entire DataFrame. To retain the old behavior, use 'frame.max(axis=0)' or just 'frame.max()'

```
    return reduction(axis=axis, out=out, **passkwargs)
```

C:\Users\PC\anaconda3\lib\site-packages\numpy\core\fromnumeric.py:84: FutureWarning: In a future version, DataFrame.min(axis=None) will return a scalar min over the entire DataFrame. To retain the old behavior, use 'frame.min(axis=0)' or just 'frame.min()'

```
    return reduction(axis=axis, out=out, **passkwargs)
```

We will split our data. 80% of our data will be train data and 20% of it will be test data.

```
In [47]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size = 0.2,random_st
```


```
In [48]: #transpose matrices
x_train = x_train.T
y_train = y_train.T
x_test = x_test.T
y_test = y_test.T
```

Let's say weight = 0.01 and bias = 0.0

```
In [49]: #initialize
def initialize(dimension):


    weight = np.full((dimension,1),0.01)
    bias = 0.0
    return weight,bias
```

## Sigmoid Function



```
In [50]: def sigmoid(z):
          y_head = 1/(1+ np.exp(-z))
          return y_head
```

## Forward and Backward Propagation



## Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$$

## Gradient Descent

### Gradient Descent

Remember that the general form of gradient descent is:

$$\begin{aligned} & \textit{Repeat} \{ \\ & \quad \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \\ & \} \end{aligned}$$

We can work out the derivative part using calculus to get:

$$\begin{aligned} & \textit{Repeat} \{ \\ & \quad \theta_j := \theta_j - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \\ & \} \end{aligned}$$

By the way in formulas;

- $h_0(x^i) = y_{\text{head}}$
- $y^i = y_{\text{train}}$
- $x^i = x_{\text{train}}$

```
In [51]: def forwardBackward(weight,bias,x_train,y_train):
# Forward

y_head = sigmoid(np.dot(weight.T,x_train) + bias)
loss = -(y_train*np.log(y_head) + (1-y_train)*np.log(1-y_head))
cost = np.sum(loss) / x_train.shape[1]

# Backward
derivative_weight = np.dot(x_train,((y_head-y_train).T))/x_train.shape[1]
derivative_bias = np.sum(y_head-y_train)/x_train.shape[1]
gradients = {"Derivative Weight" : derivative_weight, "Derivative Bias" : derivative_bias}

return cost,gradients
```



```
In [52]: def update(weight,bias,x_train,y_train,learningRate,iteration) :
    costList = []
    index = []

    #for each iteration, update weight and bias values
    for i in range(iteration):
        cost,gradients = forwardBackward(weight,bias,x_train,y_train)
        weight = weight - learningRate * gradients["Derivative Weight"]
        bias = bias - learningRate * gradients["Derivative Bias"]

        costList.append(cost)
        index.append(i)

    parameters = {"weight": weight,"bias": bias}

    print("iteration:",iteration)
    print("cost:",cost)

    plt.plot(index,costList)
    plt.xlabel("Number of Iteration")
    plt.ylabel("Cost")
    plt.show()

    return parameters, gradients
```

```
In [53]: def predict(weight,bias,x_test):
    z = np.dot(weight.T,x_test) + bias
    y_head = sigmoid(z)

    y_prediction = np.zeros((1,x_test.shape[1]))

    for i in range(y_head.shape[1]):
        if y_head[0,i] <= 0.5:
            y_prediction[0,i] = 0
        else:
            y_prediction[0,i] = 1
    return y_prediction
```

```
In [56]: def logistic_regression(x_train,y_train,x_test,y_test,learningRate,iteration):
    dimension = x_train.shape[0]
    weight,bias = initialize(dimension)

    parameters, gradients = update(weight,bias,x_train,y_train,learningRate,iteration)


    y_prediction = predict(parameters["weight"],parameters["bias"],x_test)

    print("Manual Test Accuracy: {:.2f}%".format((100 - np.mean(np.abs(y_prediction - y_test)) * 100)))
```

```
In [55]: logistic_regression(x_train,y_train,x_test,y_test,1,100)
```

iteration: 100

cost: 0.3597736123664534



Manuel Test Accuracy: 86.89%

**Manuel Test Accuracy is 86.89%**

Let's find out sklearn's score.

## Sklearn Logistic Regression

```
In [57]: accuracies = {}

lr = LogisticRegression()
lr.fit(x_train.T,y_train.T)
acc = lr.score(x_test.T,y_test.T)*100

accuracies['Logistic Regression'] = acc
print("Test Accuracy {:.2f}%".format(acc))
```


Test Accuracy 86.89%

1. **Our model works with 86.89% accuracy.**

## K-Nearest Neighbour (KNN) Classification

Let's see what will be score if we use KNN algorithm.

### KNN Algorithm



```
In [58]: # KNN Model
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors = 2) # n_neighbors means k
knn.fit(x_train.T, y_train.T)
prediction = knn.predict(x_test.T)


print("{} NN Score: {:.2f}%".format(2, knn.score(x_test.T, y_test.T)*100))
```

2 NN Score: 77.05%

```
In [59]: # try to find best k value
scoreList = []
for i in range(1,20):
    knn2 = KNeighborsClassifier(n_neighbors = i) # n_neighbors means k
    knn2.fit(x_train.T, y_train.T)
    scoreList.append(knn2.score(x_test.T, y_test.T))

plt.plot(range(1,20), scoreList)
plt.xticks(np.arange(1,20,1))
plt.xlabel("K value")
plt.ylabel("Score")
plt.show()

acc = max(scoreList)*100
accuracies['KNN'] = acc
print("Maximum KNN Score is {:.2f}%".format(acc))
```



Maximum KNN Score is 88.52%


As you can see above if we define k as 3-7-8 we will reach maximum score.

**KNN Model's Accuracy is 88.52%**

## Support Vector Machine (SVM) Algorithm

Now we will use SVM algorithm.

## Support Vector Machine Algorithm



```
In [60]: from sklearn.svm import SVC
```

```
In [61]: svm = SVC(random_state = 1)
svm.fit(x_train.T, y_train.T)

acc = svm.score(x_test.T, y_test.T)*100
accuracies['SVM'] = acc
print("Test Accuracy of SVM Algorithm: {:.2f}%".format(acc))
```

Test Accuracy of SVM Algorithm: 88.52%

**Test Accuracy of SVM Algorithm is 86.89%**

## Naive Bayes Algorithm

Naive Bayes Algorithm

The diagram shows the formula for Bayes' Theorem:  $P(c | x) = \frac{P(x | c)P(c)}{P(x)}$ . Arrows point from the terms to their labels:  $P(c | x)$  is labeled 'Posterior Probability',  $P(x | c)$  is labeled 'Likelihood',  $P(c)$  is labeled 'Class Prior Probability', and  $P(x)$  is labeled 'Predictor Prior Probability'.

$$P(c | X) = P(x_1 | c) \times P(x_2 | c) \times \cdots \times P(x_n | c) \times P(c)$$

```
In [62]: from sklearn.naive_bayes import GaussianNB
nb = GaussianNB()
nb.fit(x_train.T, y_train.T)

acc = nb.score(x_test.T, y_test.T)*100
accuracies['Naive Bayes'] = acc
print("Accuracy of Naive Bayes: {:.2f}%".format(acc))
```


Accuracy of Naive Bayes: 86.89%

**Accuracy of Naive Bayes: 86.89%**

## Decision Tree Algorithm

Decision Tree Algorithm

## Decision Tree



```
In [63]: from sklearn.tree import DecisionTreeClassifier
dtc = DecisionTreeClassifier()
dtc.fit(x_train.T, y_train.T)

acc = dtc.score(x_test.T, y_test.T)*100
accuracies['Decision Tree'] = acc
print("Decision Tree Test Accuracy {:.2f}%".format(acc))
```

Decision Tree Test Accuracy 80.33%

**Test Accuracy of Decision Tree Algorithm: 78.69%**

## Random Forest Classification

```
In [64]: # Random Forest Classification
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(n_estimators = 1000, random_state = 1)
rf.fit(x_train.T, y_train.T)

acc = rf.score(x_test.T, y_test.T)*100
accuracies['Random Forest'] = acc
print("Random Forest Algorithm Accuracy Score : {:.2f}%".format(acc))
```


Random Forest Algorithm Accuracy Score : 88.52%

## Test Accuracy of Random Forest: 88.52%

### Comparing Models

```
In [65]: colors = ["purple", "green", "orange", "magenta", "#CFC60E", "#0FBBAE"]

sns.set_style("whitegrid")
plt.figure(figsize=(16,5))
plt.yticks(np.arange(0,100,10))
plt.ylabel("Accuracy %")
plt.xlabel("Algorithms")
sns.barplot(x=list(accuracies.keys()), y=list(accuracies.values()), palette=colors)
plt.show()
```



Our models work fine but best of them are KNN and Random Forest with 88.52% of accuracy. Let's look their confusion matrixes.

### Confusion Matrix

```
In [66]: # Predicted values
y_head_lr = lr.predict(x_test.T)
knn3 = KNeighborsClassifier(n_neighbors = 3)
knn3.fit(x_train.T, y_train.T)
y_head_knn = knn3.predict(x_test.T)
y_head_svm = svm.predict(x_test.T)
y_head_nb = nb.predict(x_test.T)
y_head_dtc = dtc.predict(x_test.T)
y_head_rf = rf.predict(x_test.T)
```



```
In [67]: from sklearn.metrics import confusion_matrix

cm_lr = confusion_matrix(y_test,y_head_lr)
cm_knn = confusion_matrix(y_test,y_head_knn)
cm_svm = confusion_matrix(y_test,y_head_svm)
cm_nb = confusion_matrix(y_test,y_head_nb)
cm_dtc = confusion_matrix(y_test,y_head_dtc)
cm_rf = confusion_matrix(y_test,y_head_rf)
```

```
In [68]: plt.figure(figsize=(24,12))

plt.suptitle("Confusion Matrixes",fontsize=24)
plt.subplots_adjust(wspace = 0.4, hspace= 0.4)

plt.subplot(2,3,1)
plt.title("Logistic Regression Confusion Matrix")
sns.heatmap(cm_lr,annot=True,cmap="Blues",fmt="d",cbar=False, annot_kws={"size":

plt.subplot(2,3,2)
plt.title("K Nearest Neighbors Confusion Matrix")
sns.heatmap(cm_knn,annot=True,cmap="Blues",fmt="d",cbar=False, annot_kws={"size":

plt.subplot(2,3,3)
plt.title("Support Vector Machine Confusion Matrix")
sns.heatmap(cm_svm,annot=True,cmap="Blues",fmt="d",cbar=False, annot_kws={"size":

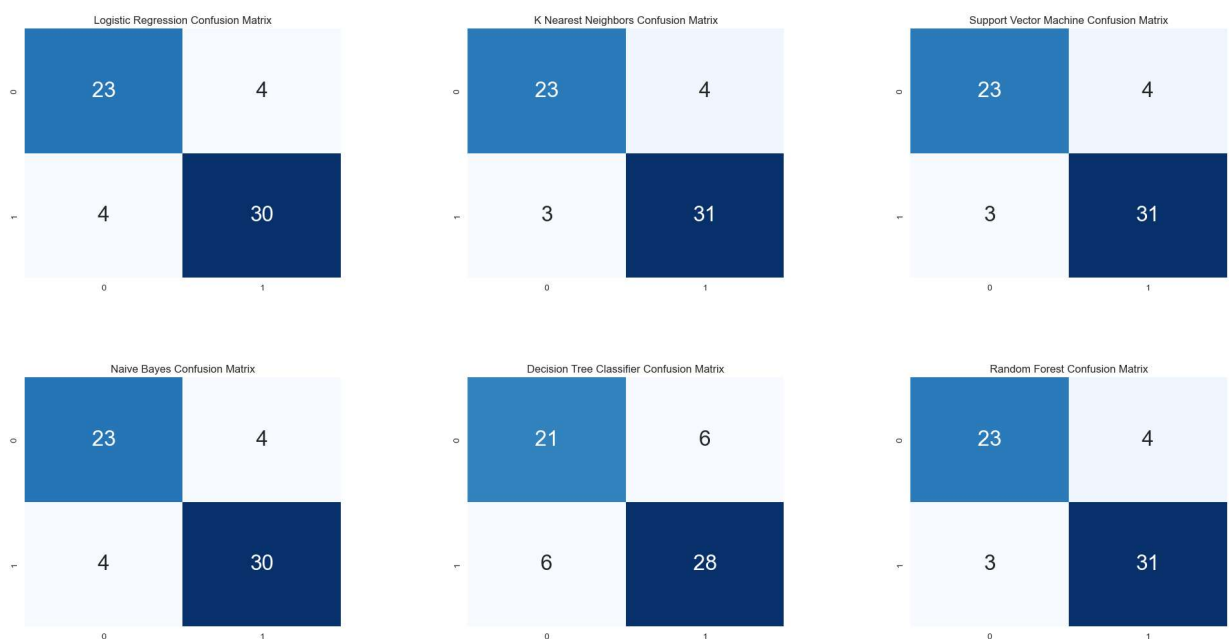
plt.subplot(2,3,4)
plt.title("Naive Bayes Confusion Matrix")
sns.heatmap(cm_nb,annot=True,cmap="Blues",fmt="d",cbar=False, annot_kws={"size":

plt.subplot(2,3,5)
plt.title("Decision Tree Classifier Confusion Matrix")
sns.heatmap(cm_dtc,annot=True,cmap="Blues",fmt="d",cbar=False, annot_kws={"size":

plt.subplot(2,3,6)
plt.title("Random Forest Confusion Matrix")
sns.heatmap(cm_rf,annot=True,cmap="Blues",fmt="d",cbar=False, annot_kws={"size":

plt.show()
```

Confusion Matrixes



\*\* Thanks for your time.\*\*

